

Dimension Modelling of Invoice Data

1. Data Profiling and Inference Process

- ▶ Overview The dataset provided contains transactional sales data with attributes related to customers, products, stores, timestamps, and dates. Based on this, a dimensional model was designed to facilitate efficient querying and analytical processing.
- ▶ Identified key attributes: Customer name, product name, store ID, invoice date, quantity, unit price, total sales, and date components.
- ▶ Examined data types and checked for inconsistencies.
- ▶ Conducted exploratory analysis to identify missing values and duplicate records.

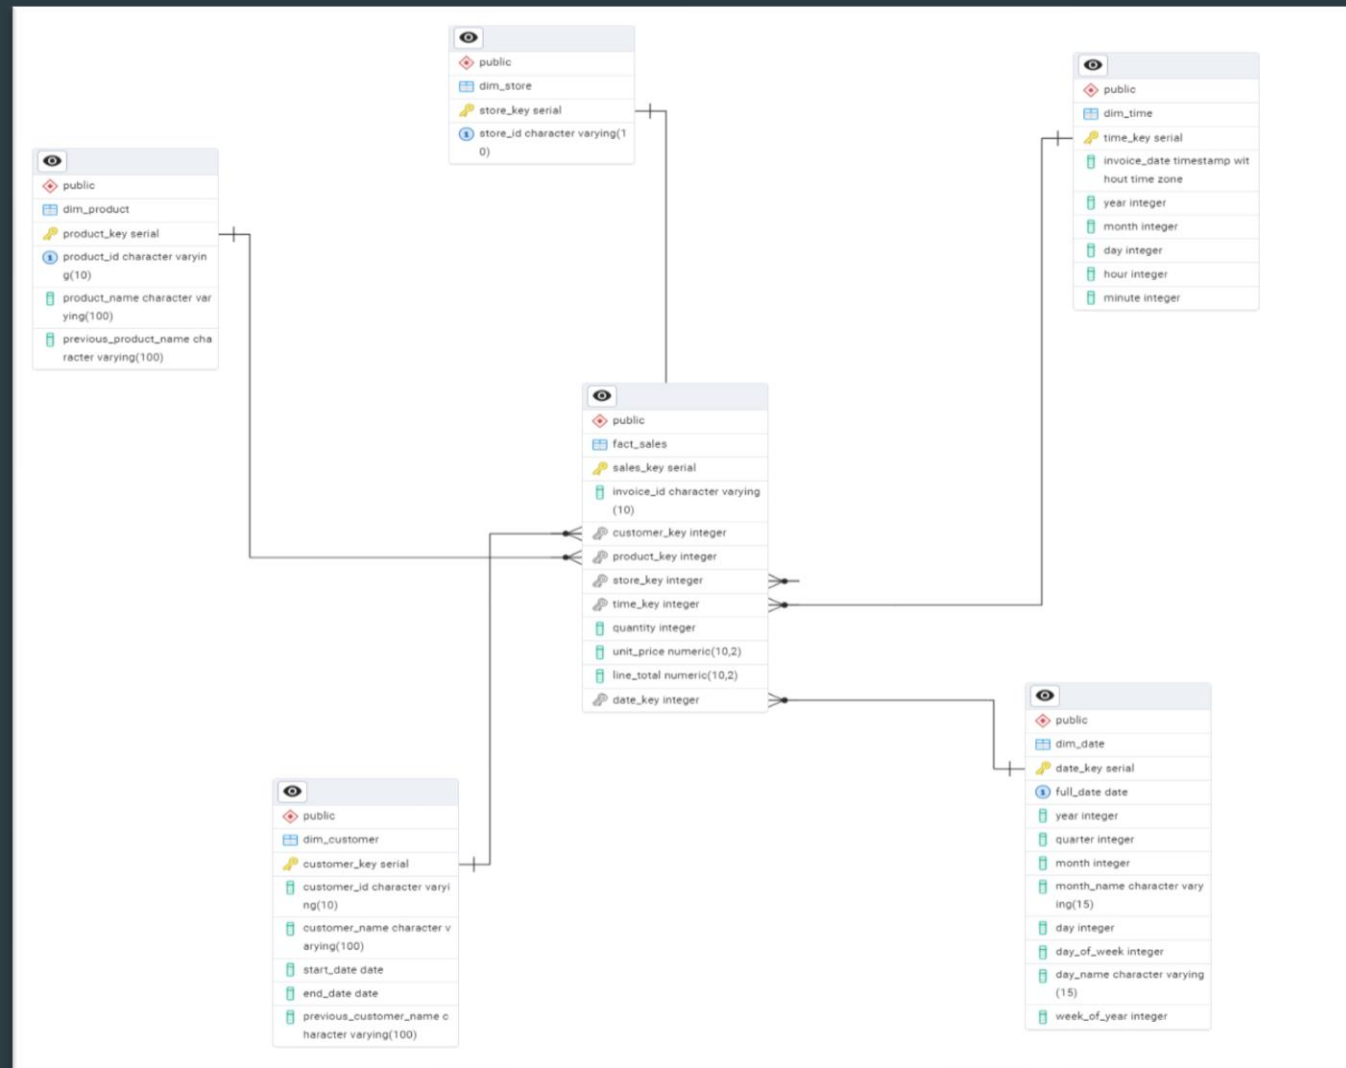
Sample Data set

Invoice_ID	Invoice_Date	Customer_ID	Customer_Name	Product_ID	Product_Name	Quantity	Unit_Price	Line_Total	Store_ID
INV001	2025-03-15 09:15:00	C001	John Smith	P1001	Wireless Mouse	2	25.00	50.00	S01
INV001	2025-03-15 09:15:00	C001	John Smith	P1002	Mechanical Keyboard	1	85.00	85.00	S01
INV002	2025-03-15 10:05:00	C002	Jane Doe	P1003	HD Monitor	1	200.00	200.00	S02
INV003	2025-03-15 11:30:00	C003	Bob Johnson	P1001	Wireless Mouse	1	25.00	25.00	S01
INV003	2025-03-15 11:30:00	C003	Bob Johnson	P1004	USB-C Hub	2	30.00	60.00	S01

Inferring Dimensions

- ▶ Customer Dimension (dim_customer): Captures customer details.
- ▶ Product Dimension (dim_product): Stores product-related attributes.
- ▶ Store Dimension (dim_store): Contains store-related identifiers.
- ▶ Time Dimension (dim_time): Extracts hierarchical time attributes for efficient time based analysis.
- ▶ Date Dimension (dim_date): Stores additional date attributes for better time-based aggregations.
- ▶ Fact Table (fact_sales): Captures transactional details linking dimensions via foreign keys.

Dimensional Model Design



Handling Slowly Changing Dimensions (SCD)

- ▶ In this schema, I implemented SCD Type 2 in the dim_customer and dim_product tables.
- ▶ This method ensures we keep track of historical changes while maintaining data integrity.
- ▶ SCD Type 2 Implementation: When a customer's name or product name changes, a new row is added with a new customer_key or product_key.
- ▶ The old record is marked with an end_date, while the new one has a start_date.
- ▶ This allows us to track changes over time and analyze past trends accurately. Example: If a customer changes their name due to marriage, we retain both old and new names to maintain accurate historical reporting.

ETL/ELT Process Flow

- ▶ To populate this schema, I followed a structured ETL process:
- ▶ Extract: We pull raw sales data, including store, customer, product, and time details.
- ▶ Transform: We standardize data formats to maintain consistency. We assign surrogate keys to ensure unique identifiers for records. Date and time are split into separate dimensions for better time-based analysis. Data is normalized into dimensions to remove redundancy and improve performance.
- ▶ Load: We insert the cleaned and structured data into fact and dimension tables. `fact_sales` stores the sales transactions with foreign keys linking to the dimension tables.

Performance Optimization

- ▶ Indexing: Primary and foreign keys indexed for fast joins.
- ▶ Partitioning: Considered for fact_sales based on time_key or date_key (not implemented due to simplicity).
- ▶ Query Optimization: Reduced sorting costs by using efficient join methods and indexing critical columns.

How This Schema Supports Business Objectives

- ▶ Track customer purchases over time.
- ▶ Analyze customer retention and churn by observing changes in dim_customer.
- ▶ Identify high-revenue products using dim_product. Compare sales performance across stores with dim_store.
- ▶ Evaluate time-based trends with dim_date and dim_time.
- ▶ Detect seasonal sales patterns. Use dim_time and dim_date for year-over-year and month-over-month comparisons.